

# CSE 587: Deep Learning for Natural Language Processing

## Lecture 3. Neural Networks and Backpropagation

Rui Zhang  
Spring 2023



PennState

# Outline - Key Concepts

NLP

N/A

ML

Neural Networks

Forward Propagation

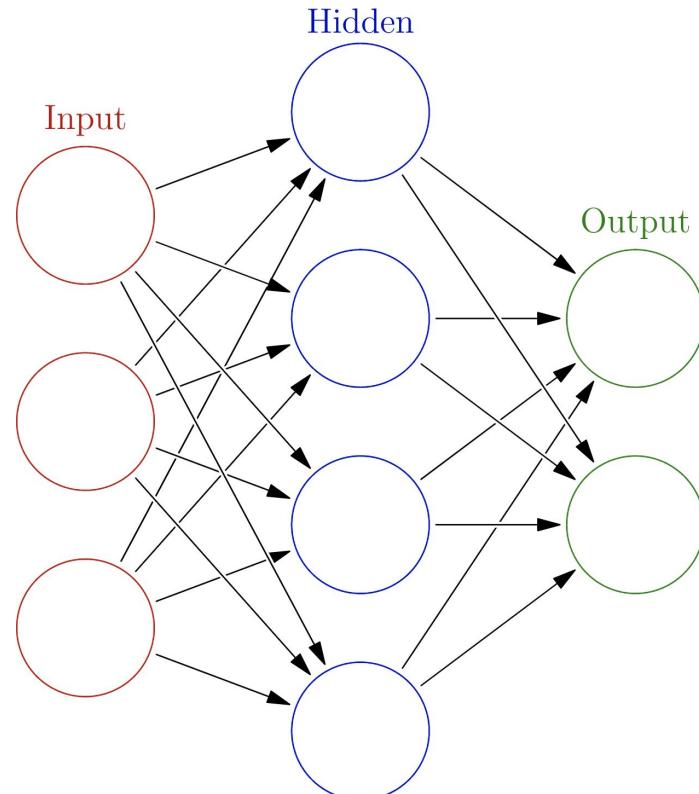
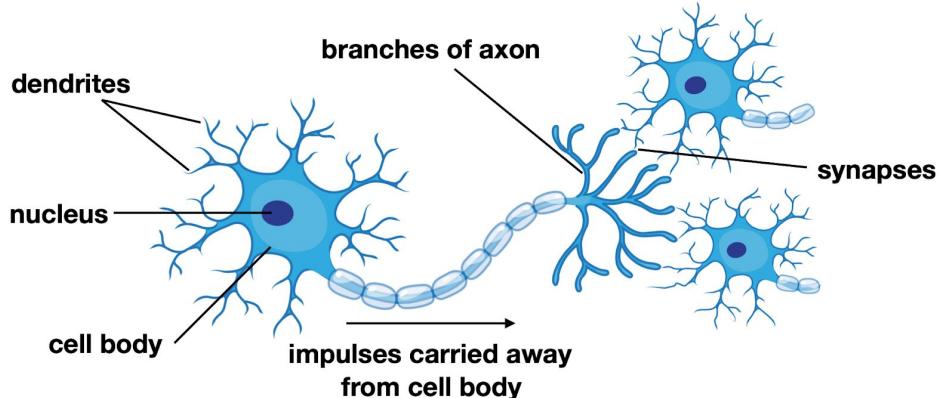
Gradient Descent

Backpropagation

Computational Graph

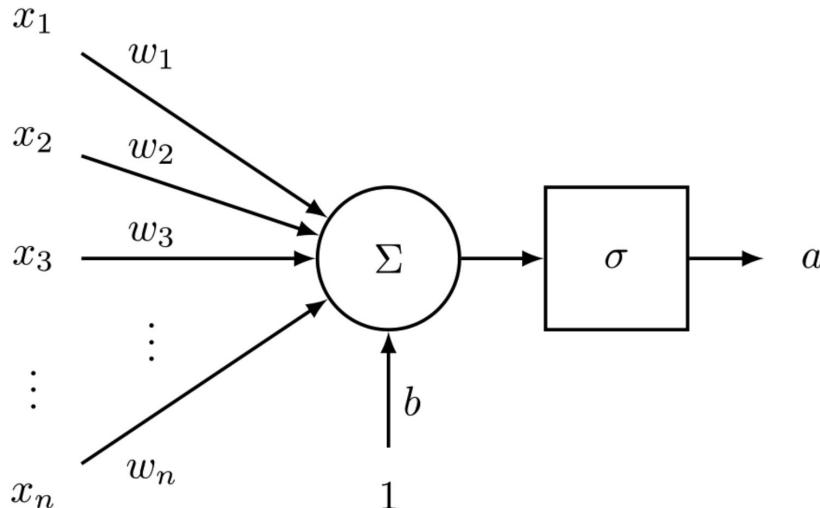
Optimizers

# Neural Networks

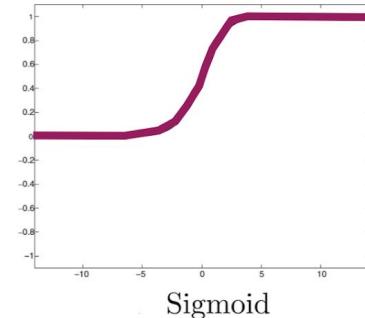


# A Neuron

A neuron is the fundamental building block of neural networks.  
e.g., a neuron with sigmoid function

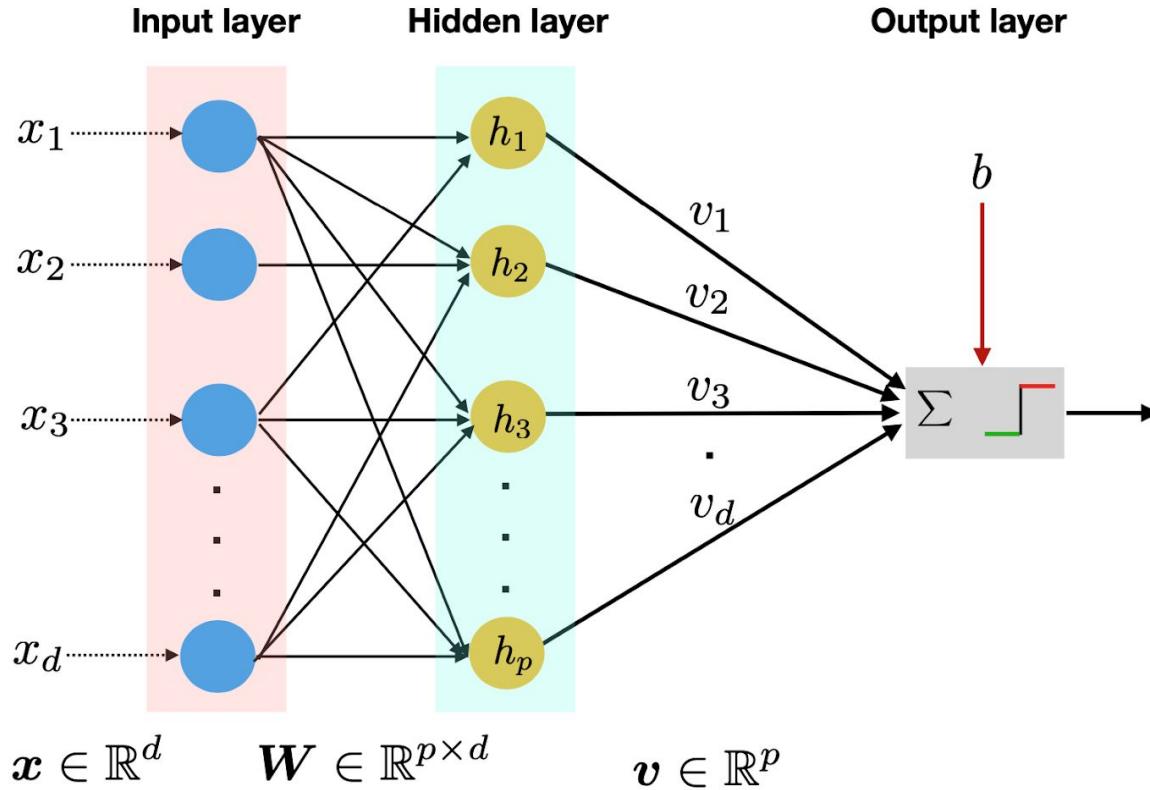


$$a = \frac{1}{1 + \exp(-(w^T x + b))}$$

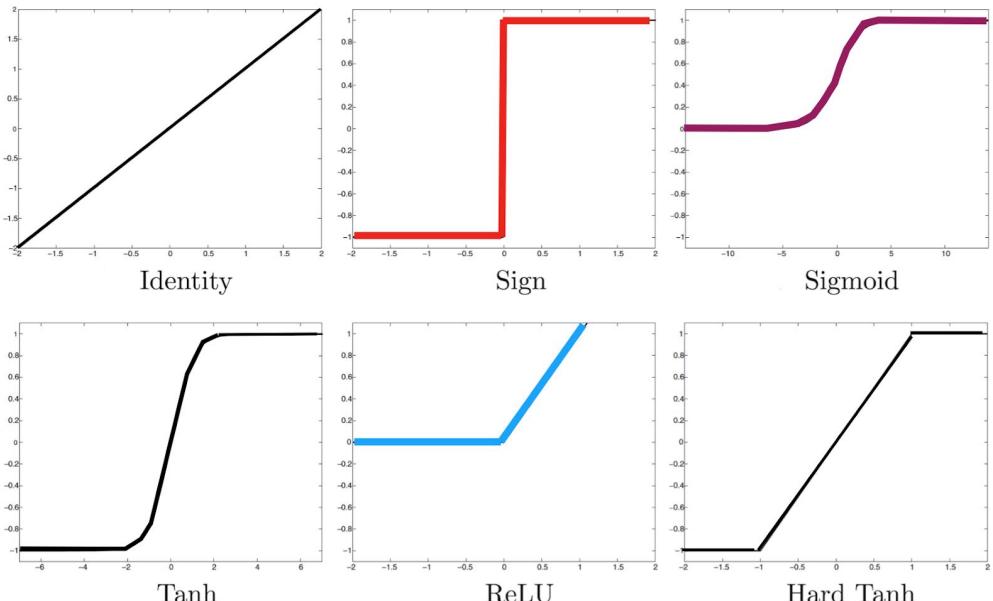


$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

# Two layer neural networks



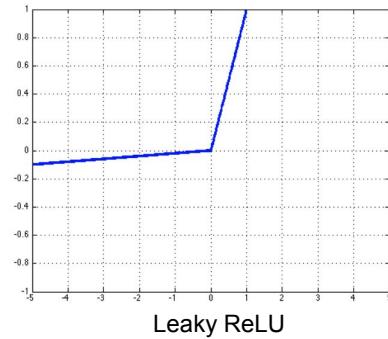
# Activation Functions



$$\sigma(z) = \text{sign}(z) \quad (\text{Sign})$$

$$\sigma(z) = \max\{0, z\} \quad (\text{Rectified Linear Unit [ReLU]})$$

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (\text{Sigmoid})$$



$$\text{leaky}(z) = \max(z, k \cdot z)$$

where  $0 < k < 1$

$$\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$$

$$\text{hardtanh}(z) = \begin{cases} -1 & : z < -1 \\ z & : -1 \leq z \leq 1 \\ 1 & : z > 1 \end{cases}$$

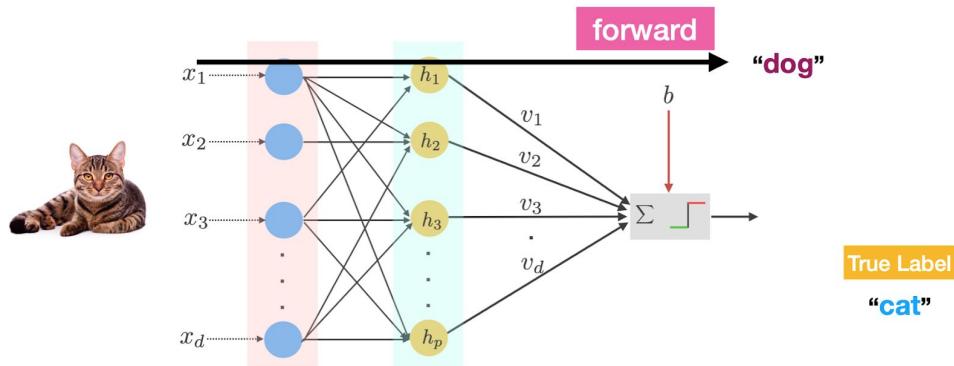
# How to train two-layer Perceptron?

**Forward propagation:**

**Backward propagation:**

# How to train two-layer Perceptron?

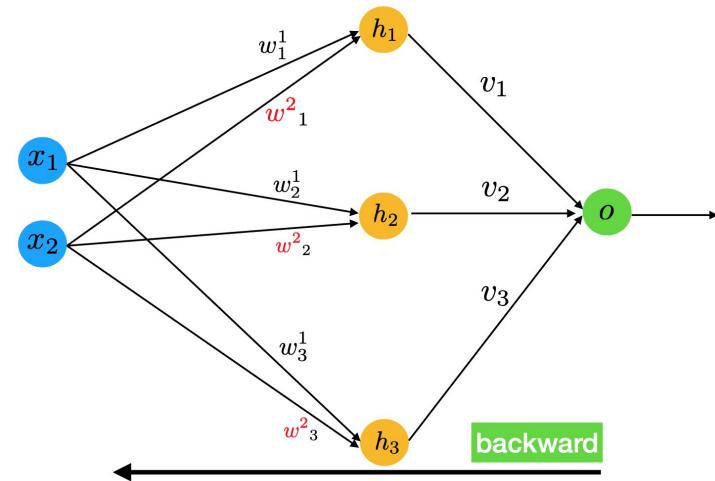
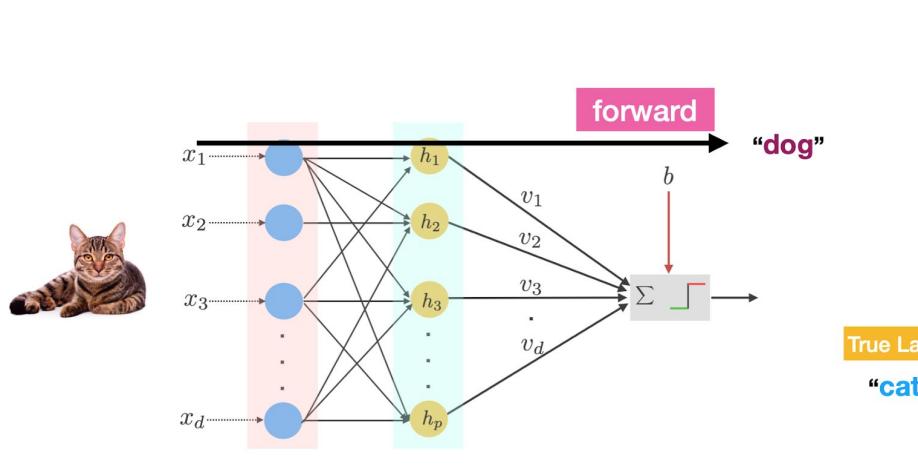
**Forward propagation:** In this phase, the inputs for a training instance are fed into the neural network. This results in a forward cascade of computations across the layers, using the current set of weights. The output is the prediction of current weights.



**Backward propagation:**

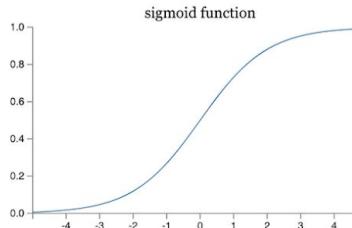
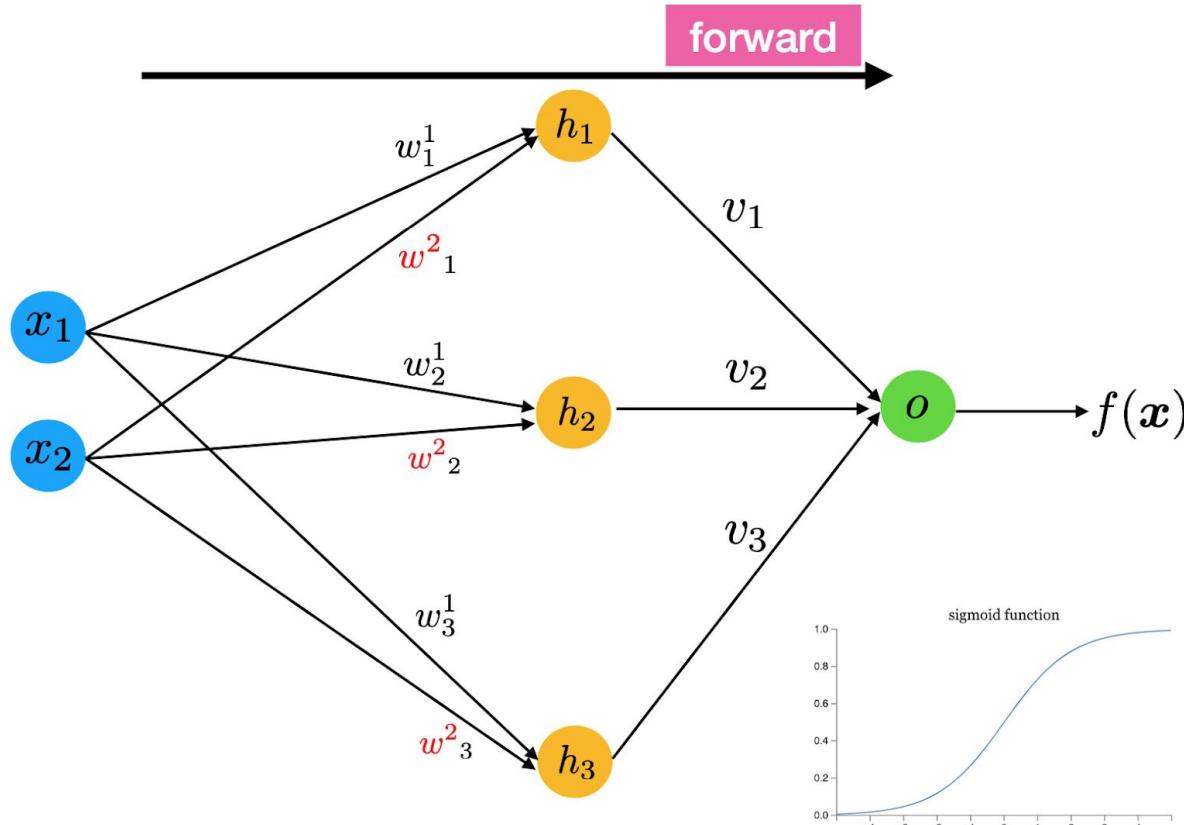
# How to train two-layer Perceptron?

**Forward propagation:** In this phase, the inputs for a training instance are fed into the neural network. This results in a forward cascade of computations across the layers, using the current set of weights. The output is the prediction of current weights.



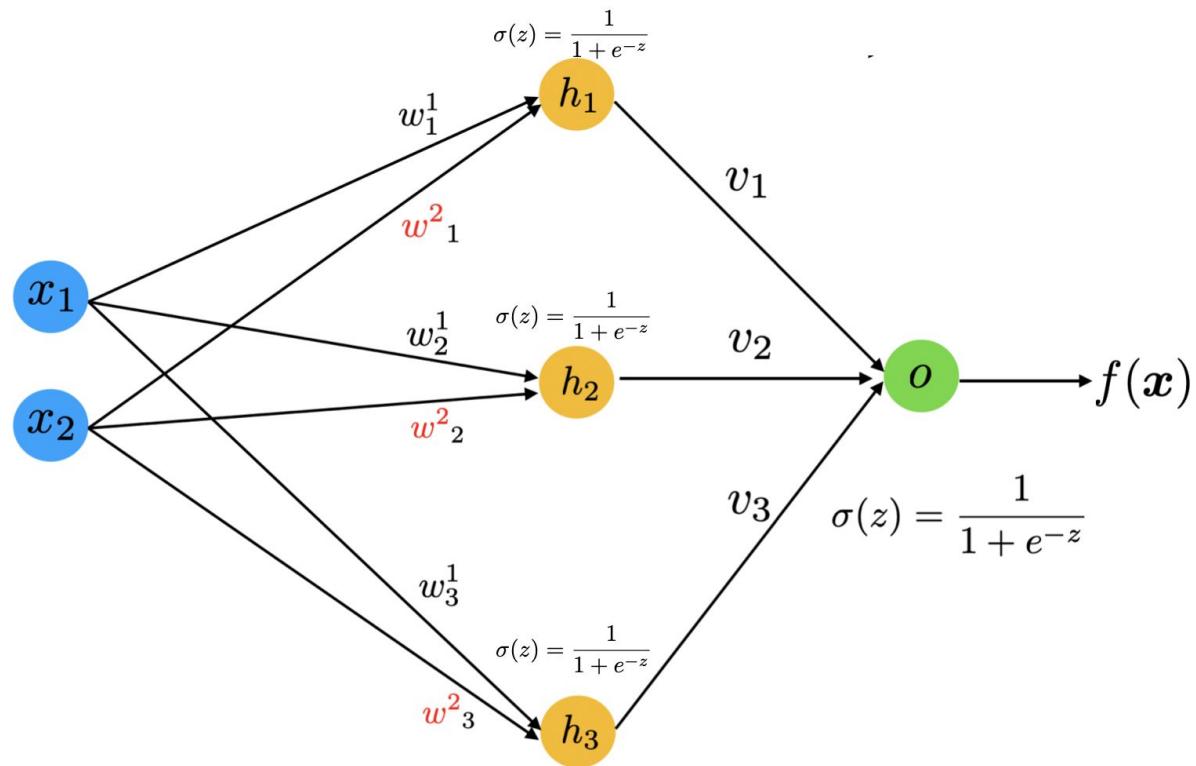
**Backward propagation:** The main goal of the backward phase is to **learn the gradient of the loss function** with respect to the different weights by **using the chain rule** of differential calculus. The weights are just adjusted based on error.

# Forward propagation

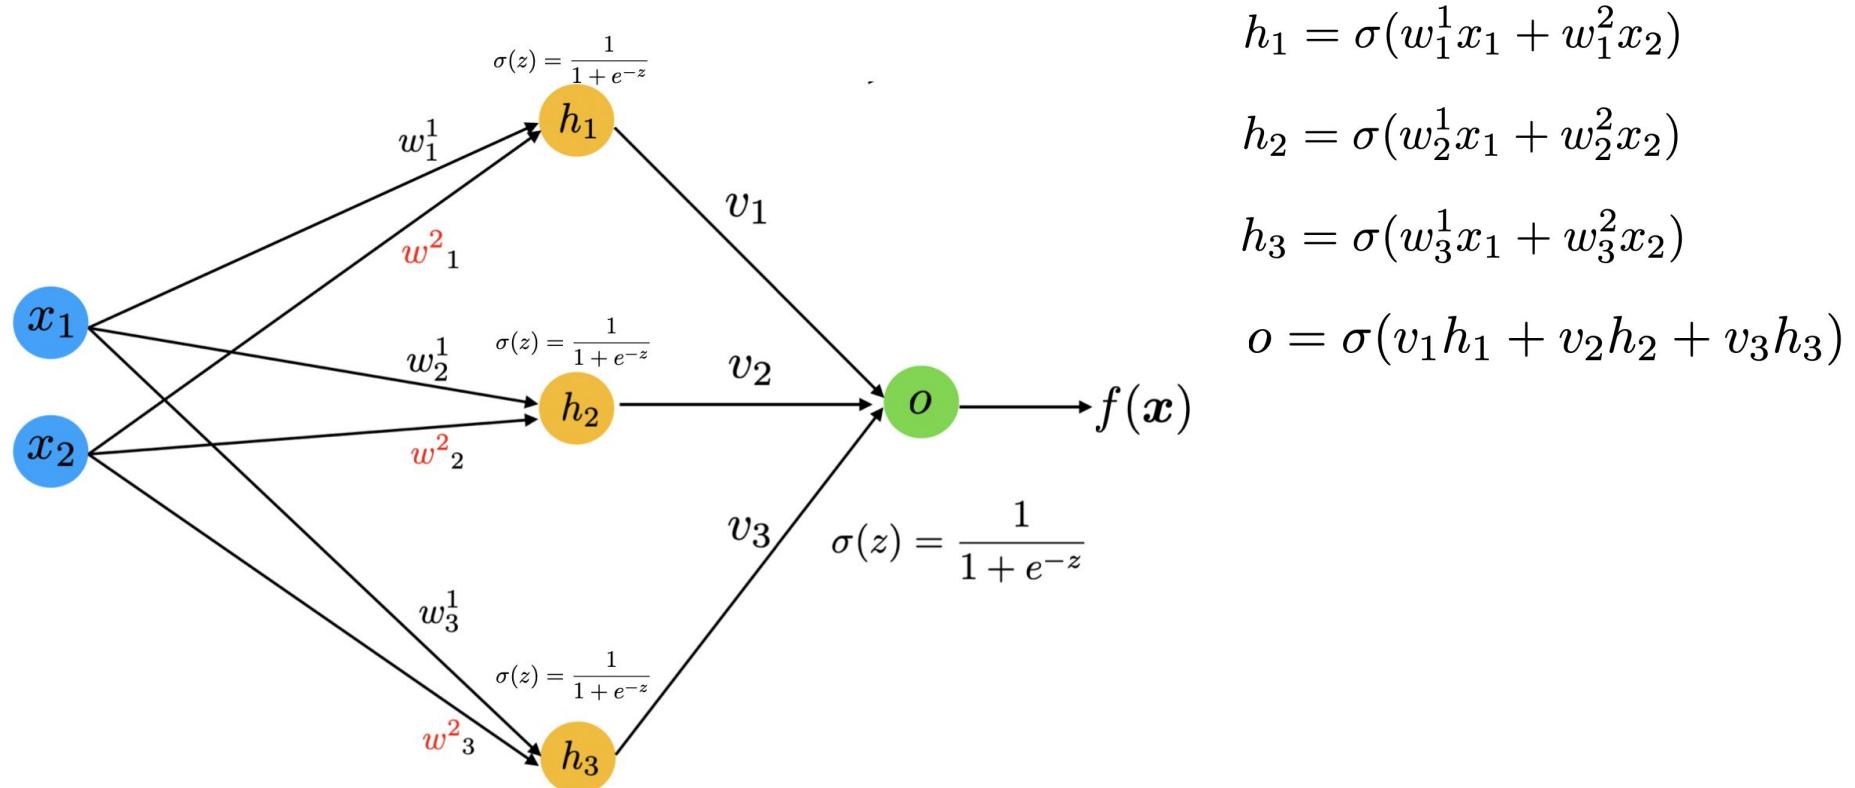


We use the **sigmoid** activation function which introduces better non-linearity:

# Forward propagation



# Forward propagation



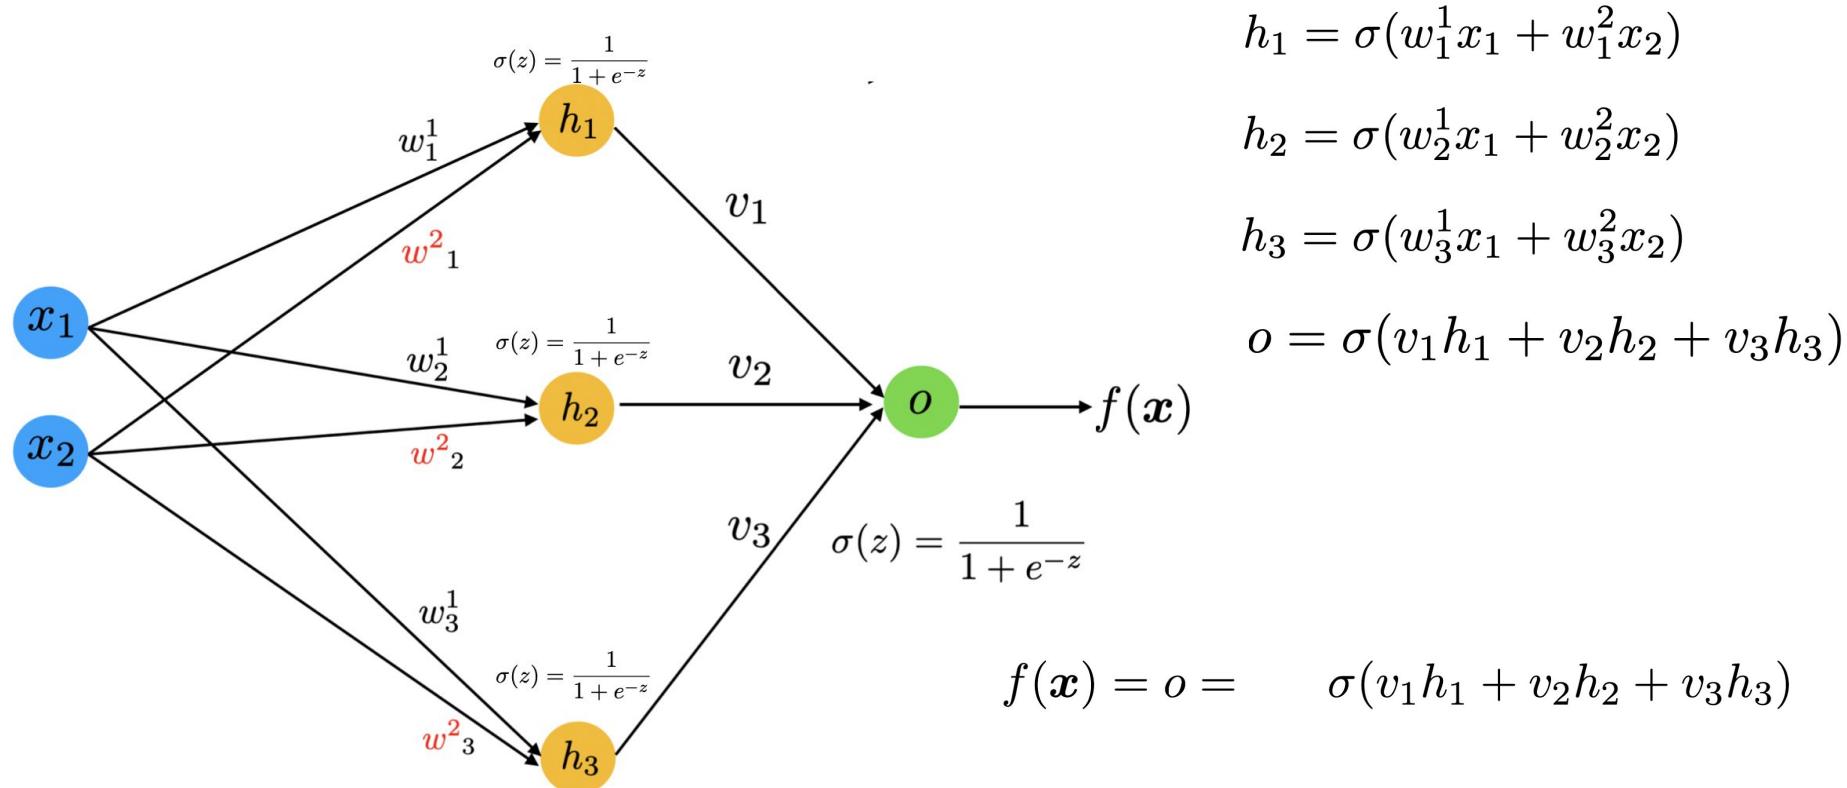
$$h_1 = \sigma(w_1^1 x_1 + w_1^2 x_2)$$

$$h_2 = \sigma(w_2^1 x_1 + w_2^2 x_2)$$

$$h_3 = \sigma(w_3^1 x_1 + w_3^2 x_2)$$

$$o = \sigma(v_1 h_1 + v_2 h_2 + v_3 h_3)$$

# Forward propagation



$$h_1 = \sigma(w_1^1 x_1 + w_1^2 x_2)$$

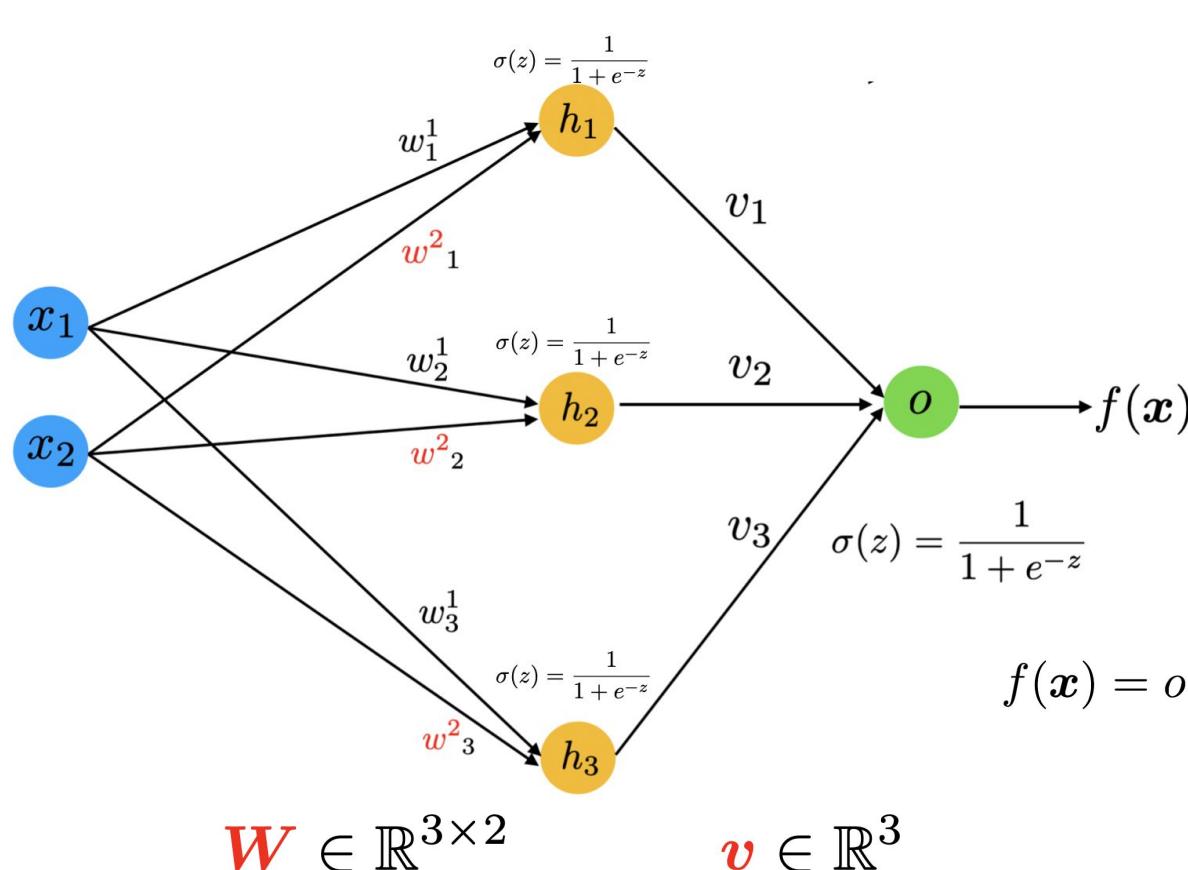
$$h_2 = \sigma(w_2^1 x_1 + w_2^2 x_2)$$

$$h_3 = \sigma(w_3^1 x_1 + w_3^2 x_2)$$

$$o = \sigma(v_1 h_1 + v_2 h_2 + v_3 h_3)$$

$$f(\mathbf{x}) = o = \sigma(v_1 h_1 + v_2 h_2 + v_3 h_3)$$

# Forward propagation - Compact Form



$$h_1 = \sigma(w_1^1 x_1 + w_1^2 x_2)$$

$$h_2 = \sigma(w_2^1 x_1 + w_2^2 x_2)$$

$$h_3 = \sigma(w_3^1 x_1 + w_3^2 x_2)$$

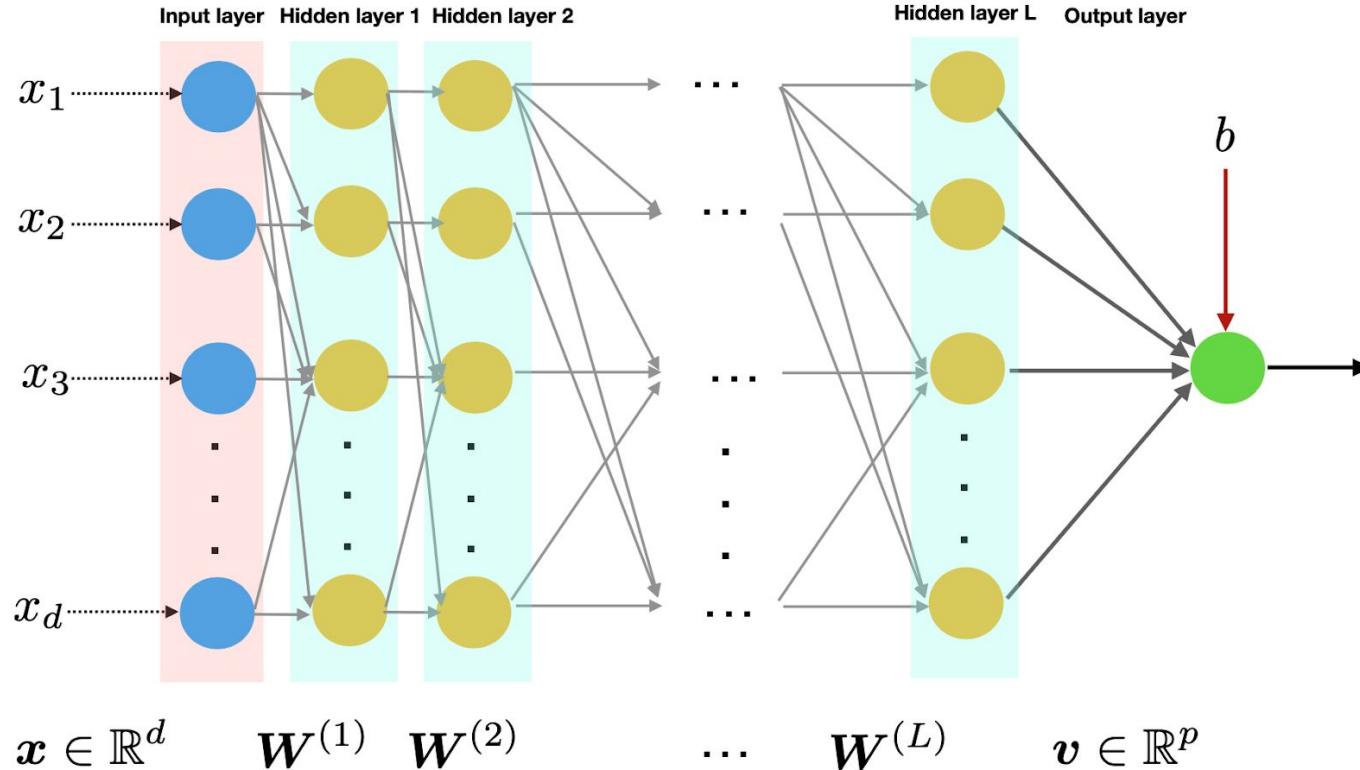
$$o = \sigma(v_1 h_1 + v_2 h_2 + v_3 h_3)$$

$$f(\mathbf{x}) = o = \sigma(v_1 h_1 + v_2 h_2 + v_3 h_3)$$

$$= \sigma \left( \mathbf{v}^\top \boldsymbol{\sigma} \left( \mathbf{W}^\top \mathbf{x} \right) \right)$$

# Beyond two layers: multilayer neural networks

Why not add more and more layers to learn more complex decision boundaries?



# Discriminative Model

The discriminative model is parameterized by  $\theta$

$$P(y|X; \theta)$$

# Discriminative Model Objective Function

The discriminative model is parameterized by  $\theta$

$$P(y|X; \theta)$$

We often use **negative log likelihood** over training data as our **objective function** or **loss function**. It is a function of  $\theta$

$$\mathcal{L}(\theta) = - \sum_{(X,y) \in \mathcal{D}_{\text{train}}} \log P(y|X; \theta)$$

# Discriminative Model Objective Function

The discriminative model is parameterized by  $\theta$

$$P(y|X; \theta)$$

We often use **negative log likelihood** over training data as our **objective function** or **loss function**. It is a function of  $\theta$

$$\mathcal{L}(\theta) = - \sum_{(X,y) \in \mathcal{D}_{\text{train}}} \log P(y|X; \theta)$$

We then optimize the parameters to minimize the loss. Better model has lower loss

$$\hat{\theta} = \arg \min_{\theta} \mathcal{L}(\theta)$$

# Optimize Objective Function by Gradient Descent

Calculate the gradient of the loss function with respect to the parameter

$$\frac{\partial \mathcal{L}(\theta)}{\partial \theta}$$

Update  $\theta$  by moving a small step in the gradient direction to decrease the loss

$$\theta_{\text{new}} \leftarrow \theta_{\text{old}} - \eta \frac{\partial \mathcal{L}(\theta)}{\partial \theta}$$

$\eta$  is the **learning rate**

# Gradient Descent

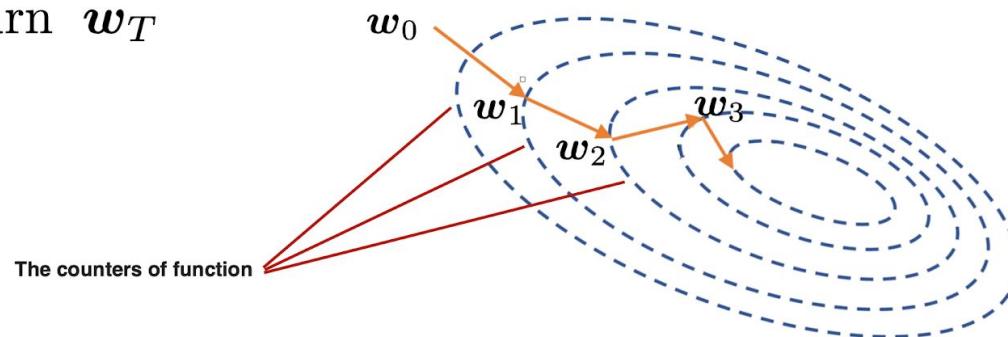
$$\mathbf{w}_0 = \mathbf{0}$$

for  $t = 1, 2, \dots, T$

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \nabla f(\mathbf{w}_t)$$

end for

return  $\mathbf{w}_T$



# Gradient Descent

The key operation in gradient descent is to calculate the gradient.

# Gradient of univariate scalar-valued functions

We are all familiar with basic calculus, and functions of the form  $f : \mathbb{R} \rightarrow \mathbb{R}$  and their derivatives:

$$f'(a) = \lim_{h \rightarrow 0} \frac{f(a + h) - f(a)}{h}.$$

The derivative of a function of a real variable measures the sensitivity to change of the function value (output value) with respect to a change in its argument (input value).

- ▶  $f(x) = x^r$ , then  $f'(x) = rx^{r-1}$ ,
- ▶  $\frac{d}{dx} e^x = e^x$ .
- ▶ Sums rule:  $(\alpha f + \beta g)' = \alpha f' + \beta g'$  for all functions  $f$  and  $g$  and all real numbers  $\alpha$  and  $\beta$
- ▶ Product rule:  $(fg)' = f'g + fg'$  for all functions  $f$  and  $g$ .
- ▶ Chain rule: If  $f(x) = h(g(x))$ , then

$$f'(x) = h'(g(x)) \cdot g'(x).$$

# Gradient of multivariate scalar-valued functions

## Definition

Let  $f : \mathcal{C} \subseteq \mathbb{R}^d \rightarrow \mathbb{R}$  be a differentiable function. Then, the gradient of  $f$  at  $\mathbf{x} \in \mathcal{C}$  is the vector in  $\mathbb{R}^d$  denoted by  $\nabla f(\mathbf{x})$  and defined by

$$f(\mathbf{x}) = f(x_1, x_2, \dots, x_n) \quad \nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1}(\mathbf{x}) \\ \vdots \\ \frac{\partial f}{\partial x_d}(\mathbf{x}) \end{bmatrix}$$

# Gradient of multivariate vector-valued functions: Jacobian Matrix

Given a function with  **$m$  outputs** and  **$n$  inputs**

$$\mathbf{f}(\mathbf{x}) = [f_1(x_1, x_2, \dots, x_n), \dots, f_m(x_1, x_2, \dots, x_n)]$$

Its Jacobian is an  **$m \times n$  matrix** of partial derivatives

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

$$\left( \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right)_{ij} = \frac{\partial f_i}{\partial x_j}$$

# Example Jacobian: Elementwise activation Function

$$\begin{aligned} \mathbf{h} = f(\mathbf{z}), \text{ what is } \frac{\partial \mathbf{h}}{\partial \mathbf{z}}? & \quad \mathbf{h}, \mathbf{z} \in \mathbb{R}^n \\ h_i = f(z_i) \end{aligned}$$

$$\left( \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \right)_{ij} = \frac{\partial h_i}{\partial z_j} = \frac{\partial}{\partial z_j} f(z_i)$$

definition of Jacobian

$$= \begin{cases} f'(z_i) & \text{if } i = j \\ 0 & \text{if otherwise} \end{cases}$$

regular 1-variable derivative

$$\frac{\partial \mathbf{h}}{\partial \mathbf{z}} = \begin{pmatrix} f'(z_1) & & 0 \\ & \ddots & \\ 0 & & f'(z_n) \end{pmatrix} = \text{diag}(\mathbf{f}'(\mathbf{z}))$$

# Other Jacobians

$$\frac{\partial}{\partial \mathbf{x}}(\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{W}$$

$$\frac{\partial}{\partial \mathbf{b}}(\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{I} \text{ (Identity matrix)}$$

$$\frac{\partial}{\partial \mathbf{u}}(\mathbf{u}^T \mathbf{h}) = \mathbf{h}^T$$

Left as exercise. Will be used later.

# Chain Rule

Chain Rule tells us how to calculate gradients of composite functions.

For composition of one-variable functions: **multiply derivatives**

$$z = 3y$$

$$y = x^2$$

$$\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx} = (3)(2x) = 6x$$

For multiple variables at once: **multiply Jacobians**

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\frac{\partial \mathbf{h}}{\partial \mathbf{x}} = \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \dots$$

# Example

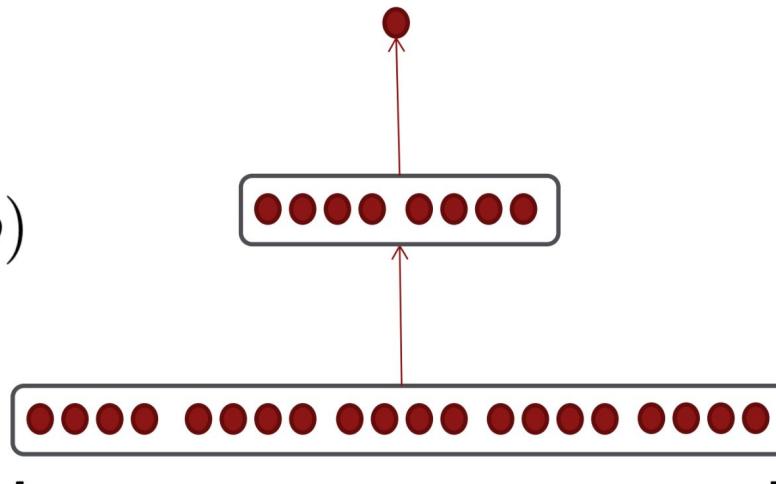
Calculate  $\frac{\partial s}{\partial b}$  for the following feed-forward neural network.

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

$\mathbf{x}$  (input)

$$\mathbf{x} = [x_{\text{museums}} \quad x_{\text{in}} \quad x_{\text{Paris}} \quad x_{\text{are}} \quad x_{\text{amazing}}]$$



# Example

Calculate  $\frac{\partial s}{\partial b}$

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{Wx} + \mathbf{b}$$

$\mathbf{x}$  (input)

$$\frac{\partial s}{\partial b} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial b}$$

# Example

Calculate  $\frac{\partial s}{\partial b}$

$$\frac{\partial}{\partial \mathbf{x}}(\mathbf{Wx} + \mathbf{b}) = \mathbf{W}$$

$$\frac{\partial}{\partial \mathbf{b}}(\mathbf{Wx} + \mathbf{b}) = \mathbf{I} \text{ (Identity matrix)}$$

$$\frac{\partial}{\partial \mathbf{u}}(\mathbf{u}^T \mathbf{h}) = \mathbf{h}^T$$

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{Wx} + \mathbf{b}$$

$\mathbf{x}$  (input)

$$\frac{\partial s}{\partial \mathbf{b}} = \frac{\partial s}{\partial \mathbf{h}} \quad \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \quad \frac{\partial \mathbf{z}}{\partial \mathbf{b}}$$

$$= \mathbf{u}^T \operatorname{diag}(f'(\mathbf{z})) \mathbf{I}$$

$$= \mathbf{u}^T \circ f'(\mathbf{z})$$

# Example

Suppose we now want to compute  $\frac{\partial s}{\partial \mathbf{W}}$

Using the chain rule again:

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)

$$\frac{\partial s}{\partial \mathbf{W}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$$

$$\frac{\partial s}{\partial \mathbf{b}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}}$$

The same! Let's avoid  
duplicated computation ...

# Example

Suppose we now want to compute  $\frac{\partial s}{\partial \mathbf{W}}$

Using the chain rule again:

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)

$$\frac{\partial s}{\partial \mathbf{W}} = \delta \boxed{\frac{\partial \mathbf{z}}{\partial \mathbf{W}}}$$

$$\frac{\partial s}{\partial \mathbf{b}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{b}} = \delta$$

$$\delta = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} = \mathbf{u}^T \circ f'(\mathbf{z})$$

# Example

How to compute  $\frac{\partial \mathbf{z}}{\partial \mathbf{W}}$ , where  $\mathbf{z} \in \mathbb{R}^n$  and  $\mathbf{W} \in \mathbb{R}^{n \times m}$

n outputs, nm inputs, by definition, the Jacobian  $\frac{\partial \mathbf{z}}{\partial \mathbf{W}}$  is  $n \times nm$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \in \mathbb{R}^m$$

$$\frac{\partial \mathbf{z}}{\partial \mathbf{W}} = \begin{bmatrix} \mathbf{x}^\top & & & \\ & \mathbf{x}^\top & & \\ & & \ddots & \\ & & & \mathbf{x}^\top \\ & & & & \mathbf{x}^\top \end{bmatrix}$$

# Example

Now we can calculate  $\frac{\partial s}{\partial W}$  of shape  $1 \times nm$

$$\frac{\partial s}{\partial W} = \delta \frac{\partial z}{\partial W} = [\delta_1 \quad \delta_2 \quad \dots \quad \delta_n] \begin{bmatrix} \mathbf{x}^\top & & & \\ & \mathbf{x}^\top & & \\ & & \ddots & \\ & & & \mathbf{x}^\top \\ & & & & \mathbf{x}^\top \end{bmatrix} = [\delta_1 \mathbf{x}^\top \quad \delta_2 \mathbf{x}^\top \quad \delta_3 \mathbf{x}^\top \quad \dots \quad \delta_n \mathbf{x}^\top]$$

# Shape Convention

The shape of the parameter is  $W \in \mathbb{R}^{n \times m}$

By definition of Jacobian, 1 output, nm inputs, so the shape of  $\frac{\partial s}{\partial W}$  is  $1 \times nm$

To make it more convenient for make gradient update, we use the  
**shape convention: the shape of the gradient is the shape of the parameters!**

So  $\frac{\partial s}{\partial W}$  is  $n$  by  $m$ :

$$\begin{bmatrix} \frac{\partial s}{\partial W_{11}} & \cdots & \frac{\partial s}{\partial W_{1m}} \\ \vdots & \ddots & \vdots \\ \frac{\partial s}{\partial W_{n1}} & \cdots & \frac{\partial s}{\partial W_{nm}} \end{bmatrix}$$

# Example

For shape convention, we reshape  $\frac{\partial s}{\partial W}$ , so it has the same shape as  $W \in \mathbb{R}^{n \times m}$

$$[\delta_1 \mathbf{x}^\top \quad \delta_2 \mathbf{x}^\top \quad \delta_3 \mathbf{x}^\top \quad \dots \quad \delta_n \mathbf{x}^\top]$$

$1 \times nm$



$$\begin{bmatrix} \delta_1 \mathbf{x}^\top \\ \delta_2 \mathbf{x}^\top \\ \delta_3 \mathbf{x}^\top \\ \vdots \\ \delta_n \mathbf{x}^\top \end{bmatrix} \quad n \times m$$

# Example

For shape convention, we reshape  $\frac{\partial s}{\partial W}$ , so it has the same shape as  $W \in \mathbb{R}^{n \times m}$

$$\begin{bmatrix} \delta_1 \mathbf{x}^\top & \delta_2 \mathbf{x}^\top & \delta_3 \mathbf{x}^\top & \dots & \delta_n \mathbf{x}^\top \end{bmatrix} \quad \xrightarrow{\hspace{1cm}} \quad \begin{bmatrix} \delta_1 \mathbf{x}^\top \\ \delta_2 \mathbf{x}^\top \\ \delta_3 \mathbf{x}^\top \\ \vdots \\ \delta_n \mathbf{x}^\top \end{bmatrix} \quad n \times m$$

$1 \times nm$

A more direct way: compute the **outer product** of  $\boldsymbol{\delta}$  and  $\mathbf{x}$   
and this directly makes the shape  $n \times m$

$$\begin{aligned} \frac{\partial s}{\partial W} &= \boldsymbol{\delta}^T \quad \mathbf{x}^T \\ [n \times m] \quad [n \times 1][1 \times m] \\ &= \begin{bmatrix} \delta_1 \\ \vdots \\ \delta_n \end{bmatrix} [x_1, \dots, x_m] = \begin{bmatrix} \delta_1 x_1 & \dots & \delta_1 x_m \\ \vdots & \ddots & \vdots \\ \delta_n x_1 & \dots & \delta_n x_m \end{bmatrix} \end{aligned}$$

# Backpropagation

We use **gradient descent** to optimize the parameters in neural networks.

The key question: How can we efficiently calculate the gradients in neural networks (or any computational graph)?

The key intuition: View neural networks (or any computational graph) as compositions of functions and use **chain rules** to calculate the gradient.

This is **Backpropagation = Gradient Descent + Chain Rule**

# Computation Graph

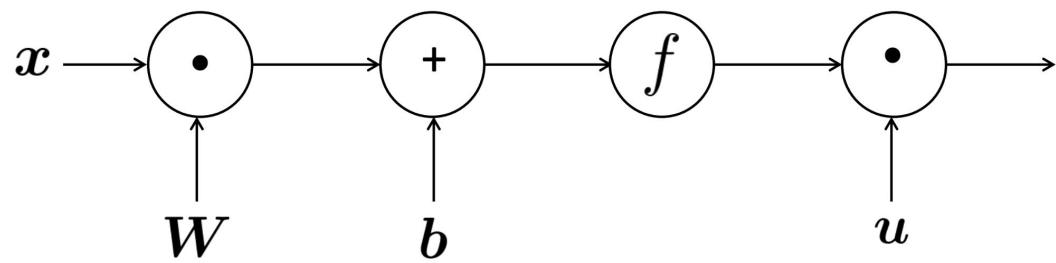
In software library, neural networks are implemented as Computation Graph (CG).

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)



# Forward Propagation in CG

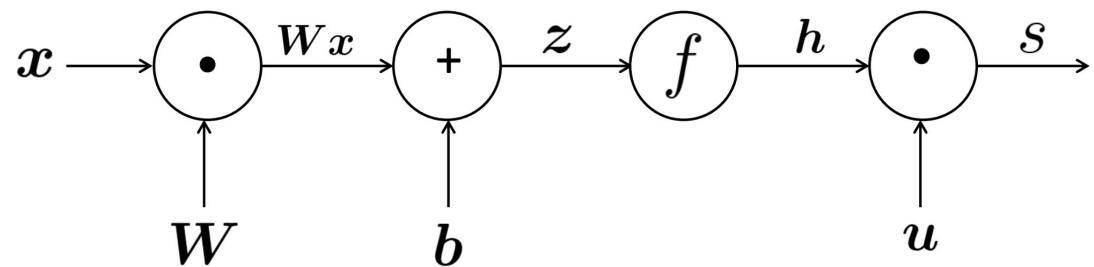
In software library, neural networks are implemented as Computation Graph (CG).

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)



# Backpropagation in CG

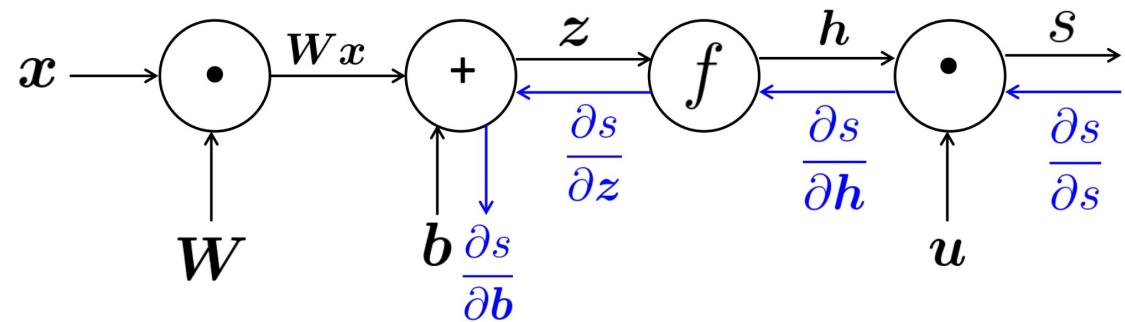
In software library, neural networks are implemented as Computation Graph (CG).

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)



# Backpropagation avoids duplicated computation

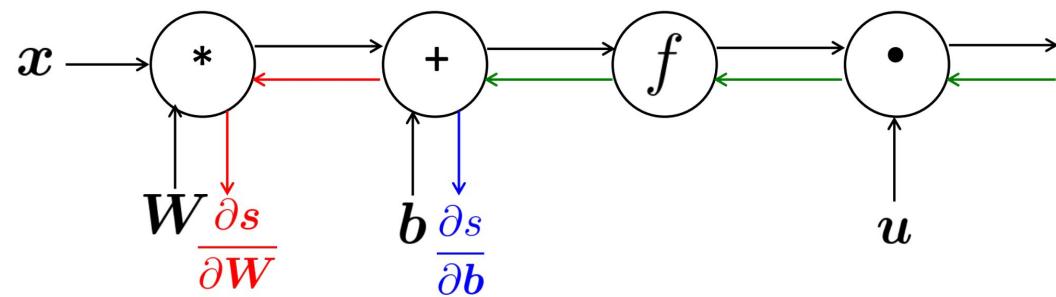
Backpropagation follows the computation graph in reverse order, so you avoid duplicated computation.

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$\mathbf{x}$  (input)



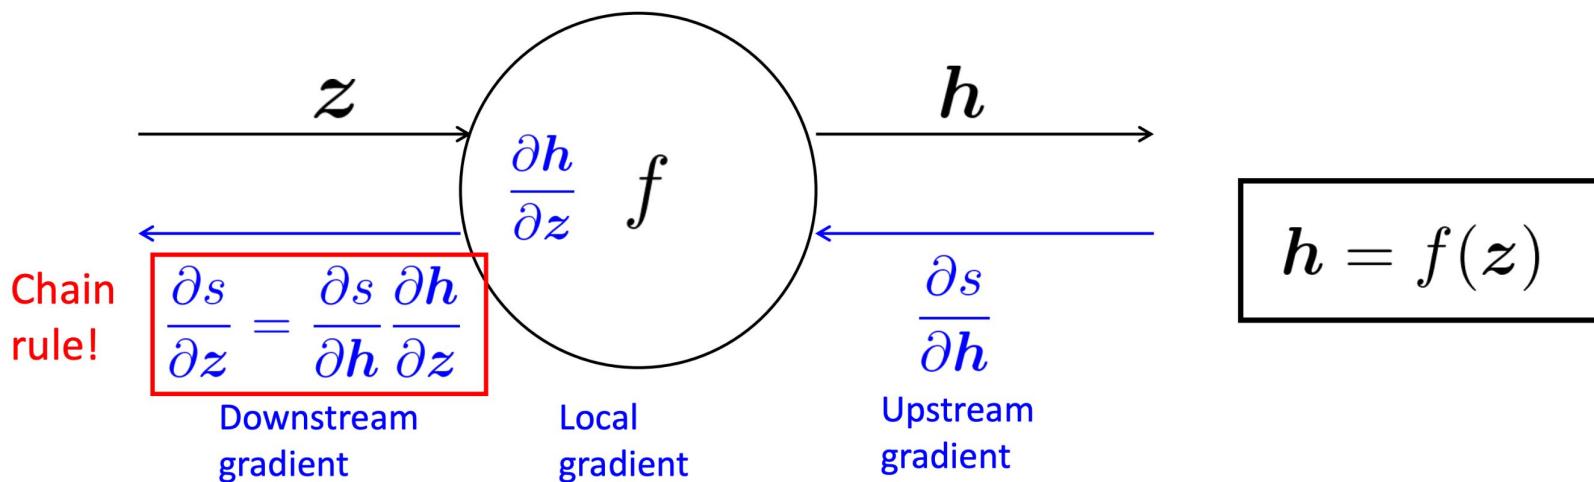
# Backpropagation: Single Node

Node receives an **upstream gradient**.

Each node has a **local gradient**.

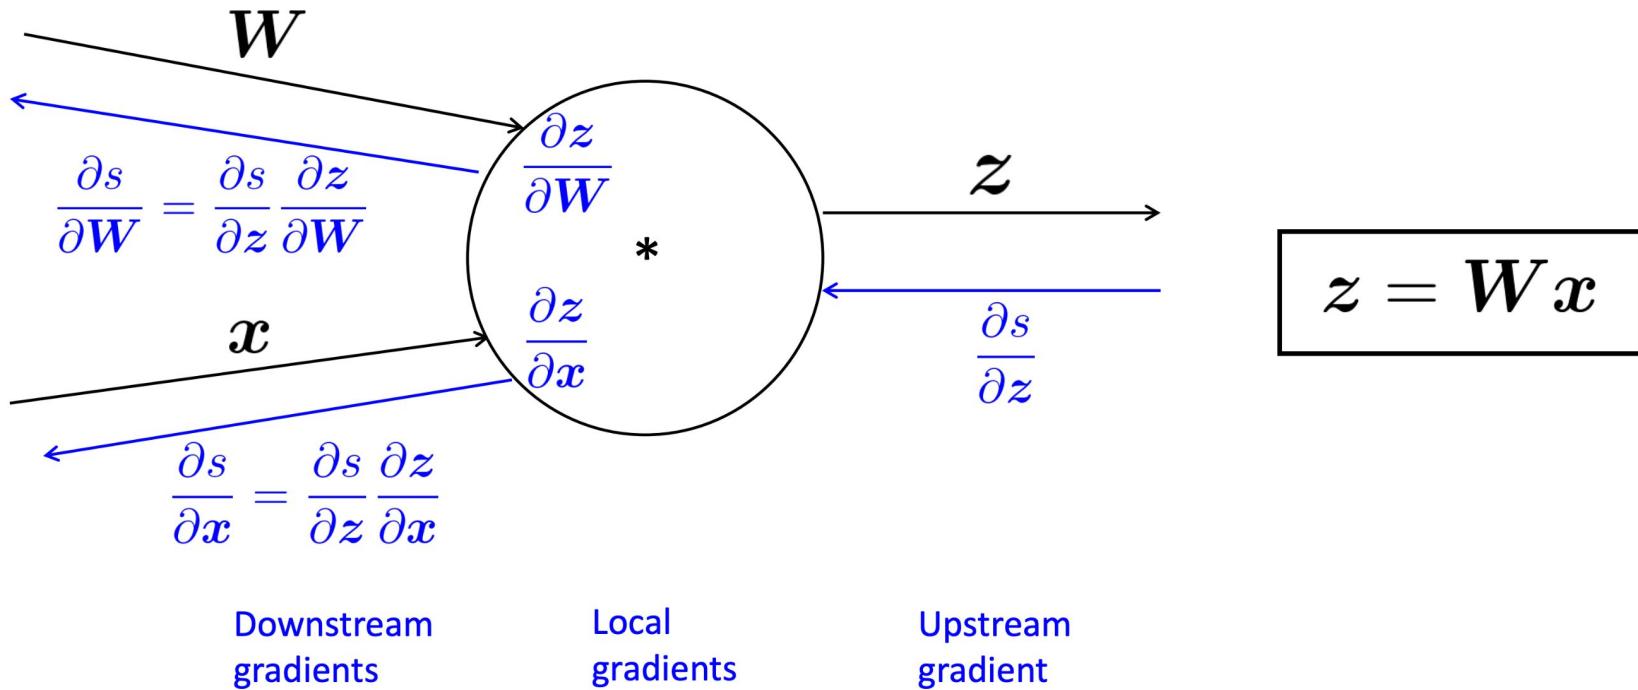
Goal is to pass on the correct **downstream gradient**.

By Chain Rule: **downstream gradient = upstream gradient x local gradient.**



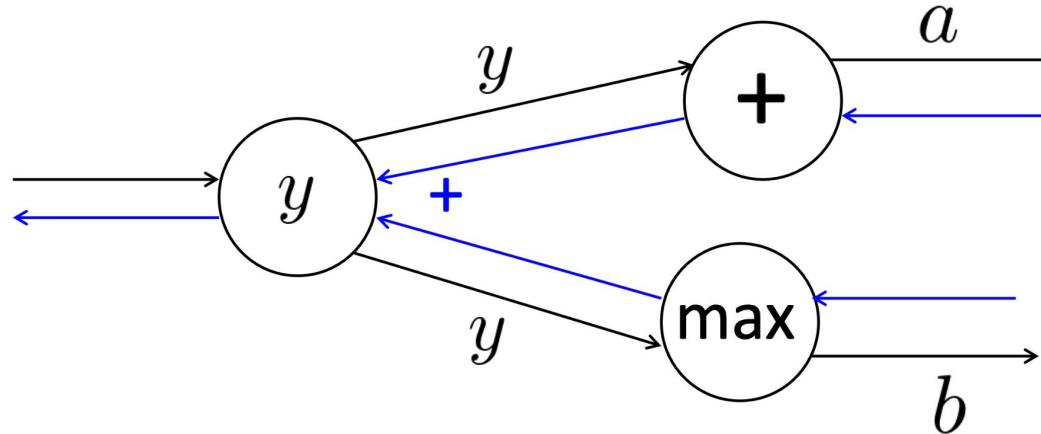
# Backpropagation: Multiple Inputs

With multiple inputs, each input has a local gradient.



# Backpropagation: Multiple Outputs

If a node has multiple outgoing edges during forward propagation, we sum the gradients during the backpropagation.



$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial a} \frac{\partial a}{\partial y} + \frac{\partial f}{\partial b} \frac{\partial b}{\partial y}$$

# Example

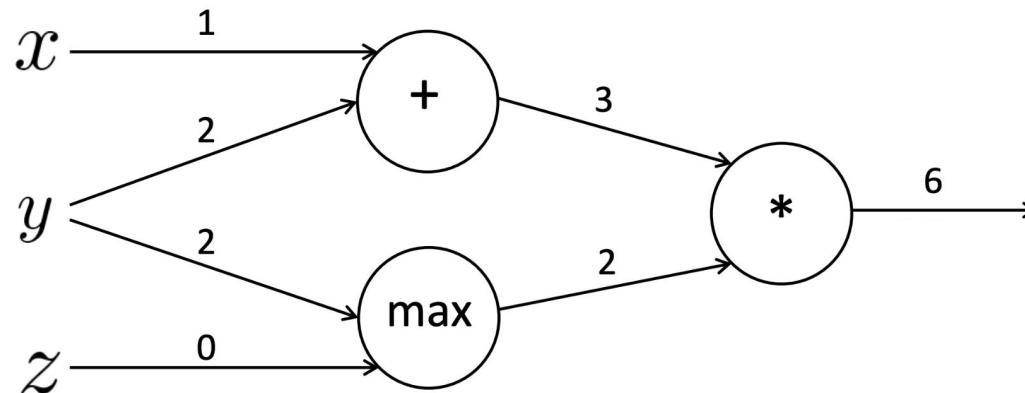
$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$



# Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

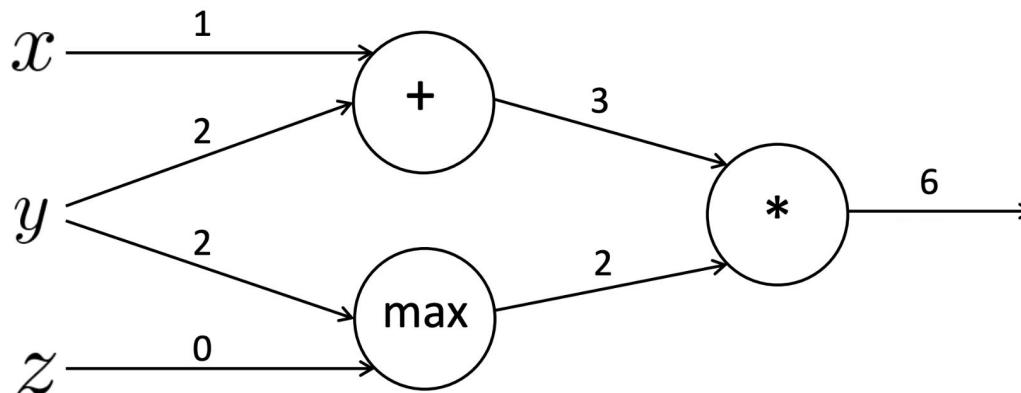
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

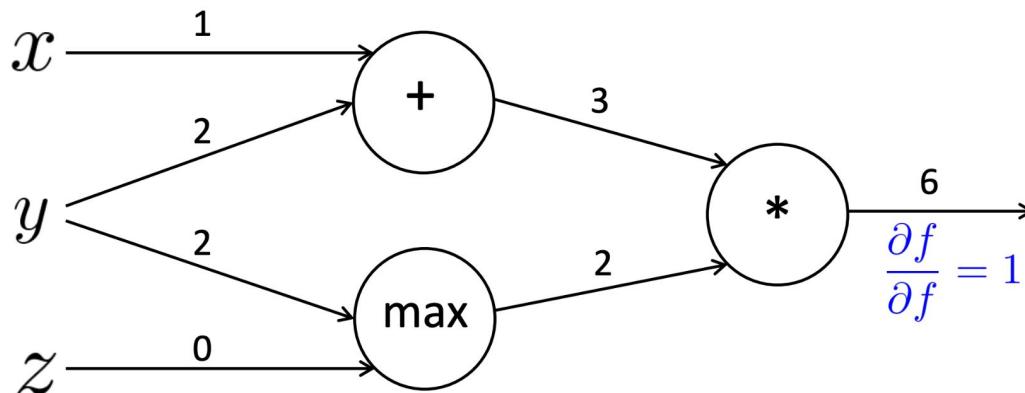
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

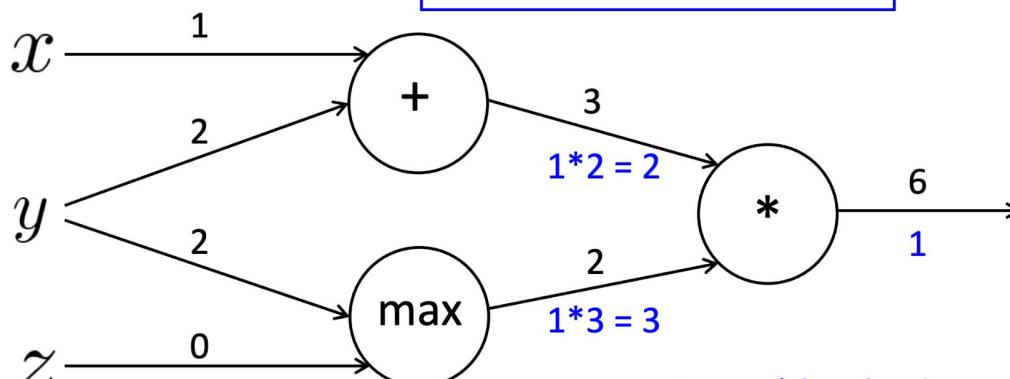
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\boxed{\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3}$$



# Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

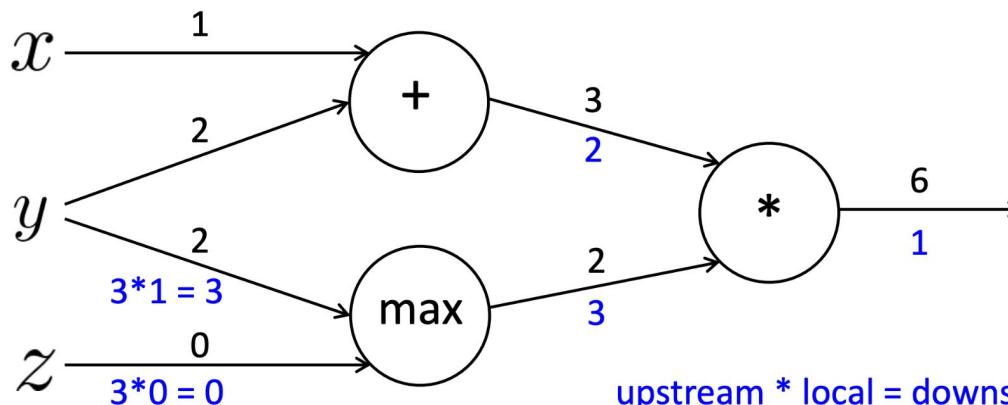
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



upstream \* local = downstream

# Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

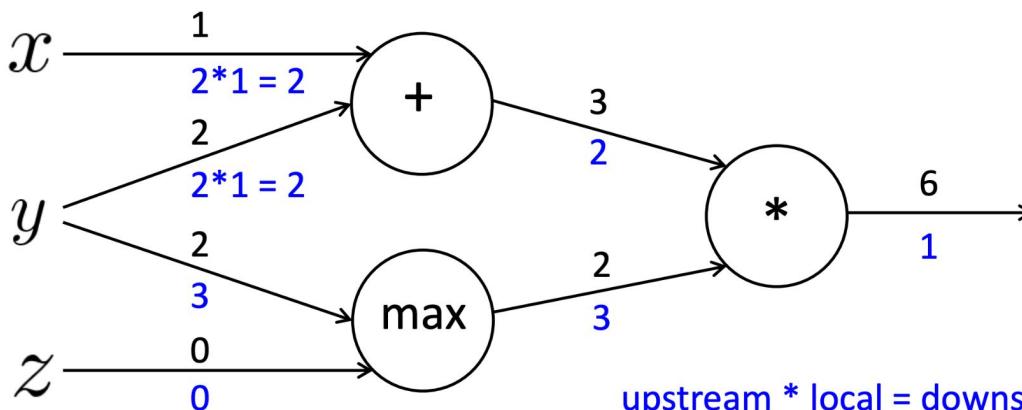
$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# Example

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

Forward prop steps

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

$$\frac{\partial f}{\partial x} = 2$$

$$\frac{\partial f}{\partial y} = 3 + 2 = 5$$

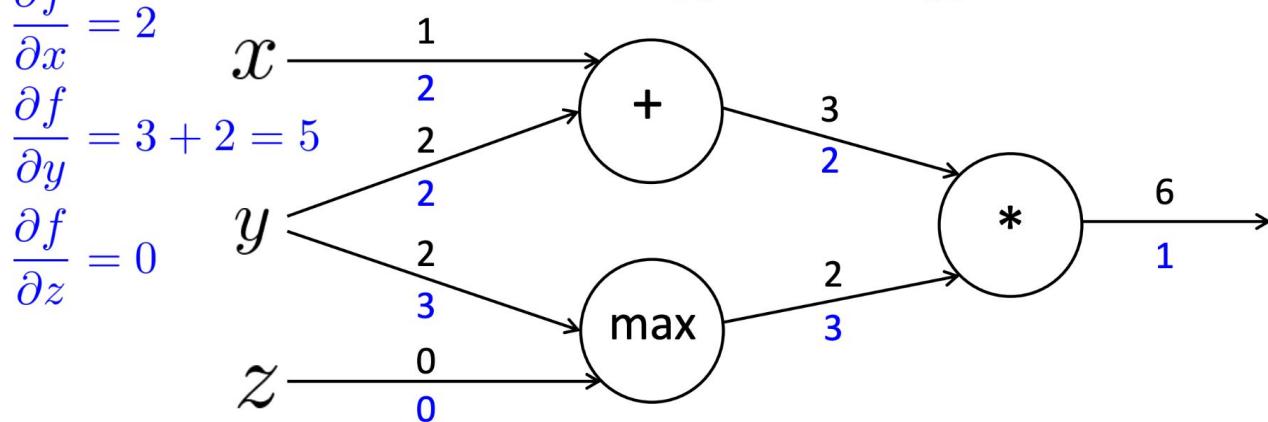
$$\frac{\partial f}{\partial z} = 0$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

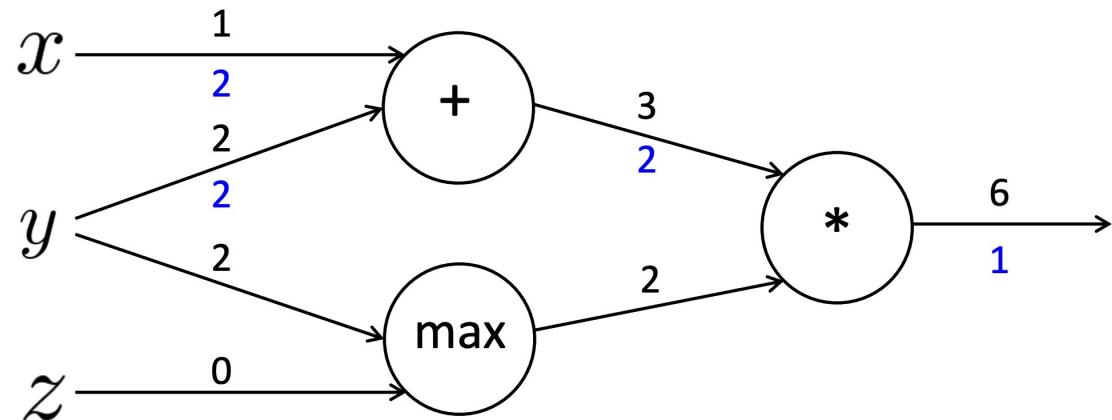
$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# Node Intuitions

+ “distributes” the upstream gradient to each summand

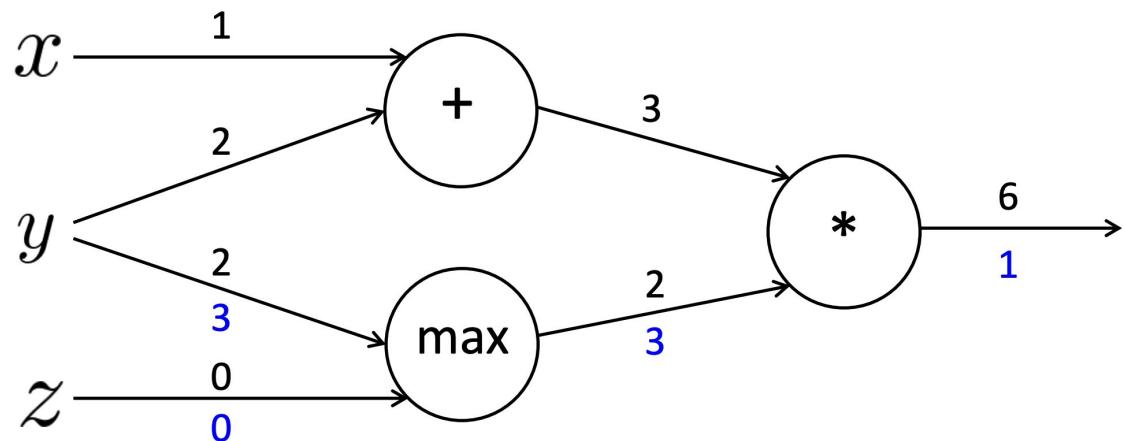
$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$



# Node Intuitions

- + “distributes” the upstream gradient to each summand
- max “routes” the upstream gradient

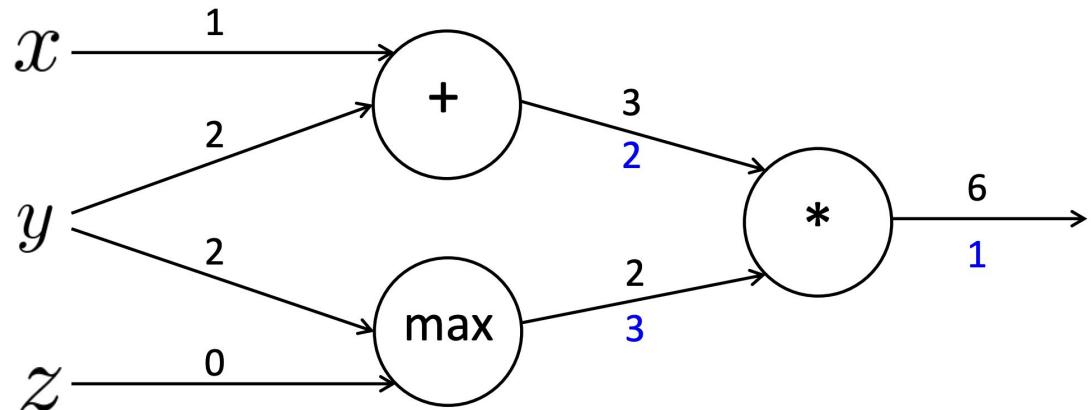
$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$



# Node Intuitions

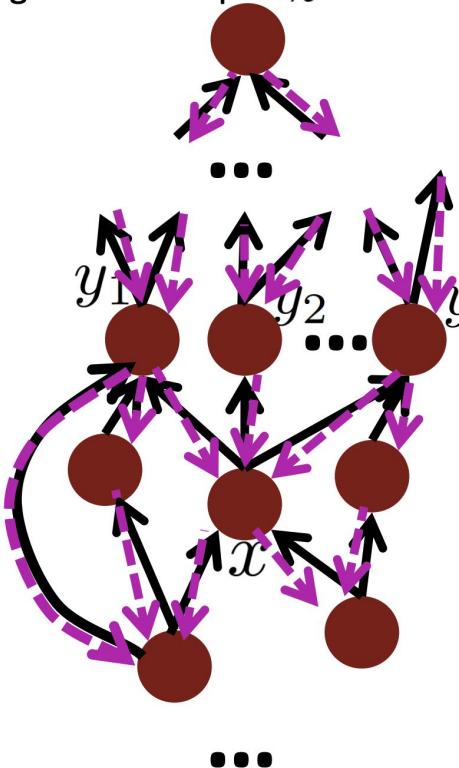
- + “distributes” the upstream gradient to each summand
- max “routes” the upstream gradient
- \* “switches” the upstream gradient

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$



# Backpropagation in General Computation Graph

Single scalar output  $z$



- F-prop: visit nodes in topological sort order - Compute value of node given predecessors.
- B-prop:
  - initialize output gradient = 1
  - visit nodes in reverse order: Compute gradient wrt each node using gradient wrt successors

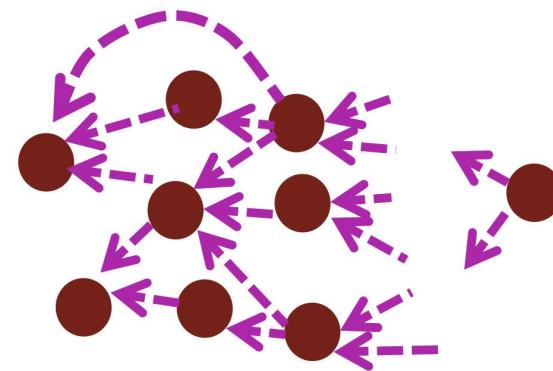
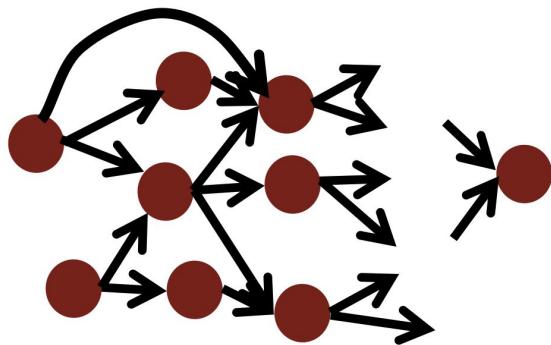
$\{y_1, y_2, \dots, y_n\}$  = successors of  $x$

$$\frac{\partial z}{\partial x} = \sum_{i=1}^n \frac{\partial z}{\partial y_i} \frac{\partial y_i}{\partial x}$$

# AutoDiff: Automatic Differentiation

The gradient computation can be automatically inferred from the symbolic expression of the fprop.

Each node type needs to know how to compute its output and how to compute the gradient wrt its inputs given the gradient wrt its output.



# Check Your Implementation with Numeric Gradient

Following the definition of gradient:

$$f'(x) \approx \frac{f(x + h) - f(x - h)}{2h}$$

Use a very small  $h$  to compute this value numerically.

Useful for checking your implementation.

# Optimizers

Once we calculate the gradient with respect to the objective function, we optimize objective function by gradient descent.

The very basic one is standard gradient descent. But there are many other choices of optimizers:

- Batch Gradient Descent
- Stochastic Gradient Descent (SGD)
- Mini-batch Gradient Descent
- SGD with Momentum
- Nesterov accelerated gradient
- Adagrad
- Adadelta
- RMSProp
- **Adam; AdamW (most commonly used)**
- Adafactor

# Batch Gradient Descent

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta)$$

New Parameters      Old Parameters      Learning Rate      Gradient

```
for i in range(nb_epochs):
    params_grad = evaluate_gradient(loss_function, data, params)
    params = params - learning_rate * params_grad
```

# Mini-batch Gradient Descent

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i:i+n)}; y^{(i:i+n)})$$



Gradient for a batch of training examples; batch size n

```
for i in range(nb_epochs):
    np.random.shuffle(data)
    for batch in get_batches(data, batch_size=50):
        params_grad = evaluate_gradient(loss_function, batch, params)
        params = params - learning_rate * params_grad
```

# Stochastic Gradient Descent (SGD)

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i)}; y^{(i)})$$



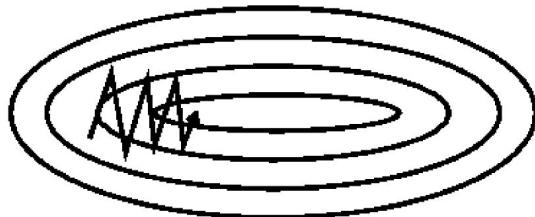
Gradient for each training example

```
for i in range(nb_epochs):
    np.random.shuffle(data)
    for example in data:
        params_grad = evaluate_gradient(loss_function, example, params)
        params = params - learning_rate * params_grad
```

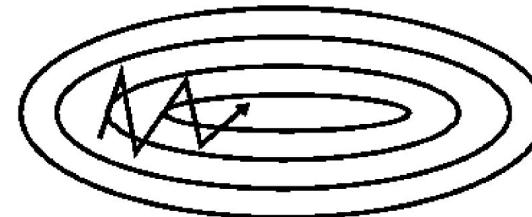
# SGD with Momentum

Momentum parameter, usually 0.9 or a similar value

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$$
$$\theta = \theta - v_t$$



(a) SGD without momentum



(b) SGD with momentum

Figure 2: Source: Genevieve B. Orr

# Nesterov Accelerated Gradient (NAG)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$

$$\theta = \theta - v_t$$



calculating the gradient not w.r.t. to our current parameters  $\theta$  but w.r.t. the approximate future position of our parameters

# Nesterov Accelerated Gradient (NAG)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$

$$\theta = \theta - v_t$$

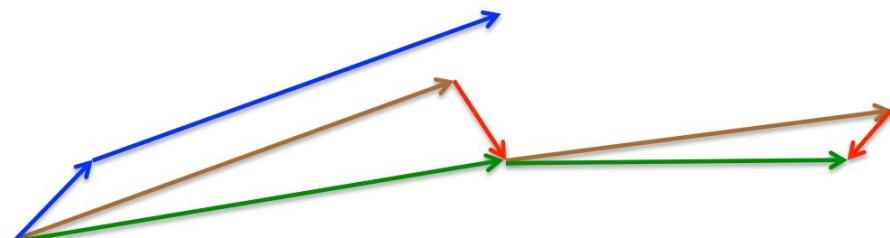
## Momentum

- first computes the current gradient (small blue vector)
- then takes a big jump in the direction of the updated accumulated gradient (big blue vector)

## NAG

- first makes a big jump in the direction of the previous accumulated gradient (brown vector)
- measures the gradient and then makes a correction (red vector)
- results in the complete NAG update (green vector).

calculating the gradient not w.r.t. to our current parameters  $\theta$  but w.r.t. the approximate future position of our parameters



# Adagrad

Adagrad adapts the learning rate to the parameters: **larger updates for infrequent parameters, smaller updates for frequent parameters.**

It uses a different learning rate for every parameter  $\theta_i$  at every time step  $t$ ,

$$g_{t,i} = \nabla_{\theta_t} J(\theta_{t,i})$$

$$G_t \in \mathbb{R}^{d \times d}$$

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}$$

$$\epsilon$$

A diagonal matrix where each diagonal element is the sum of the squares of the gradients w.r.t.  $\theta_i$  up to time step  $t$

smoothing term that avoids division by zero

# Adagrad

Adagrad adapts the learning rate to the parameters: **larger updates for infrequent parameters, smaller updates for frequent parameters.**

It uses a different learning rate for every parameter  $\theta_i$  at every time step  $t$ ,

$$g_{t,i} = \nabla_{\theta_t} J(\theta_{t,i})$$

$$G_t \in \mathbb{R}^{d \times d}$$

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}$$

$$\epsilon$$

A diagonal matrix where each diagonal element is the sum of the squares of the gradients w.r.t.  $\theta_i$  up to time step  $t$

smoothing term that avoids division by zero

Problem: learning rate continuously decreases, and training can stall -- fixed by using **rolling average** in AdaDelta, RMSProp, Adam.

# Adaptive Moment Estimation (Adam)

$m_t$  and  $v_t$  are estimates of the first moment (the mean) and the second moment (the uncentered variance) of the gradients respectively, hence the name of the method.

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad \xleftarrow{\hspace{1cm}} \text{like momentum}$$
$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \quad \xleftarrow{\hspace{1cm}} \text{rolling average!}$$

# Adaptive Moment Estimation (Adam)

$m_t$  and  $v_t$  are estimates of the first moment (the mean) and the second moment (the uncentered variance) of the gradients respectively, hence the name of the method.

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad \longleftarrow \quad \text{like momentum}$$
$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \quad \longleftarrow \quad \text{rolling average!}$$

As  $m_t$  and  $v_t$  are initialized as vectors of 0's, the authors of Adam observe that they are biased towards zero, especially during the initial time steps, and especially when the decay rates are small (i.e.  $\beta_1$  and  $\beta_2$  are close to 1). They counteract these biases by computing bias-corrected first and second moment estimates,

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$
$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

Then update the parameters

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

# Training Tricks

**Shuffling the Training Data Every Epoch:** avoid providing the training examples in a meaningful order to our model as this may bias the optimization algorithm.

**Regularization:** L2 regularization adds L2 norm of parameters to the loss function.

**Dropout:** randomly zero-out nodes in the hidden layer with probability  $p$  at training time only

**Learning Rate Decay:** gradually reduce learning rate as training continues

## Early stopping

- You should thus always monitor error on a validation set during training and stop (with some patience) if your validation error does not improve enough.
- Use Patience

# Parameter Initialization

Very Important for neural networks to achieve good performance.

**Uniform Initialization:** Initialize weights in some range, such as [-0.1, 0.1] for example

**Xavier Initialization:**  $n(l)$  is the number of input units to  $W$  (fan-in) and  $n(l+1)$  is the number of output units from  $W$  (fan-out).

$$W \sim U\left[-\sqrt{\frac{6}{n^{(l)} + n^{(l+1)}}}, \sqrt{\frac{6}{n^{(l)} + n^{(l+1)}}}\right]$$